



// Engineering Meetups

Core.async: Deep Dive




Andrey Listopadov,
Clojure Engineer at Health Samurai

About me

- Clojure engineer at Health Samurai
- Love asynchronous programming
- Interested in various programming languages
 - Helping with development of the Fennel programming language

My personal site



GitHub / GitLab: [@andreyorst](#)

What'll this talk be about

- Asynchronous programming
 - When/Why?
 - History of async in Clojure
- `clojure.core.async`
 - Introduction to core.async
 - Implementation of go-threads
 - Implementation of channels
 - Timers
 - Event Loop/Scheduler
- Portability
- Problems

Asynchronous programming

- Async in a language either
 - Exists as a part of the language
 - Attached by some means
- Async in other languages
 - BEAM - actor model
 - Go - CSP
 - .NET - async/await
 - Lua - coroutines
- Async in JVM
 - Green Threads (before JDK 1.1)
 - ...nothing?
 - Virtual Threads, since JDK21

Parallelism VS Async – what's the difference?

- Parallelism
 - Useful mostly for compute tasks
 - Parallel execution
 - Limited by core count and their abilities
- Async
 - Useful for IO-bound tasks
 - Cooperative execution
 - Doesn't depend on core count (depends)
 - Limited by the acceptable time of reaction to events

Async in Clojure

- Agents
 - Agent is a “place” for storing modifiable data as agent state
 - Agent’s state is modified by sending tasks
 - Tasks - pure functions that accept current state and return a new state
 - Tasks are serialized in a queue, everything is thread-safe
- Third-party alternatives
 - Manifold – event-driven
 - Promesa – async/await
 - Probably some others...

Agents: Problems

- Tasks should remain short
- Blocking an agent can:
 - Exhaust agent's thread pool - no more tasks will run
 - Cause OOME by growing queue indefinitely, since there's no backpressure
- Agents have two separate thread pools:
 - Compute: unbound, for compute tasks
 - IO: bound, for IO tasks
- Error handling
 - In a complex network of messages it's often hard to find where exactly the error happened
 - Stacktraces are hard to understand
 - Managing agent's working state adds a bit of complexity to the code
- Can't be ported to all other supported runtimes!

`clojure.core.async`

- Universal async model for all supported runtimes
- CSP-based (Communicating Sequential Processes)
- Platform-independent asynchronous thread implementation
- Channels as synchronisation points, and process communication primitive

Channels

- A simple queue*
- Unbuffered:
 - Supports backpressure by default
- Buffered
 - Backpressure once buffer can't accept more
- Different buffering strategies:
 - Fixed buffer: backpressure, if full
 - Sliding Window: **N** most recent messages
 - Dropping buffer: **N** oldest messages
 - Custom:
 - Drop random messages?
 - Drop **N** messages then enable backpressure?
 - Other crazy ideas

With a channel you can!

- Put data in! >!
- Get data out! <!
- Put data in, or block the thread!! >!!
- Get data out, or block the thread!! <!!

Blocking ops

- Blocks the thread currently executed in
- Better to avoid in asynchronous threads

```
1. (def ch (chan 10))  
2. (>!! ch "Health Samurai") ;; puts "Health Samurai" into channel's buffer  
3. (println (<!! ch)) ;; Prints Health Samurai to stdout
```

Parking ops

- Do not block anything
- Can be used only inside asynchronous threads
 - It's a compile-time error to use parking ops outside of asynchronous threads

```
1. (def ch (chan))  
2. (go  
3.   (println (<! ch)))  
4. (go  
5.   (>! ch "Health Samurai"))
```

Asynchronous threads

- Platform independent implementation
- Lightweight
- Fast to create
- Behaves as a channel

go threads

- Go threads – are not threads at all
- **go** – is a macro, that transforms code inside to behave like a thread

```
1. (def ch1 (chan))  
2. (def ch2 (chan))  
3. (go  
4.   (let [data (<! ch1)]  
5.     (>! ch2 (foo data))))
```

go's output

```
1.  (let [c (chan 1)
2.        captured-bindings (getThreadBindingFrame)
3.        f (fn state-machine
4.              ([] (aset-all! (AtomicReferenceArray. (int 6)) 0 state-machine 1 1))
5.              ([state]
6.                  (let [old-frame (getThreadBindingFrame)
7.                        ret (try
8.                                (resetThreadBindingFrame (aget-object state 3)))
9.                                (loop []
10.                                    (let [result
11.                                         (case (int (aget-object state 1))
12.                                             1 (take! state 2 ch1)
13.                                             2 (let [data (foo (aget-object state 2))]
14.                                                 (put! state 3 ch2 data))
15.                                             3 (return-chan state (aget-object state 2))))
16.                                         (if (identical? result :recur) (recur) result)))
17.                                    (catch Throwable ex
18.                                        (aset-all! state 2 ex)
19.                                        (if (seq (aget-object state 4))
20.                                            (aset-all! state 1 (first (aget-object state 4)))
21.                                            (throw ex)))
22.                                        :recur)
23.                                    (finally
24.                                        (aset-object state 3 (getThreadBindingFrame)) (resetThreadBindingFrame old-frame)))
25.                                    (if (identical? ret :recur)
26.                                        (recur state)
27.                                        ret))))
28.              state (aset-all! (f) USER-START-IDX c BINDINGS-IDX captured-bindings)]
29.              (run-state-machine-wrapped state)
30.              c)
```

```
(def ch1 (chan))
(def ch2 (chan))
(go
  (let [data (<! ch1])
    (>! ch2 (foo data))))
```

go's output

```
1.  (let [c (chan 1) ;; ❶ creating a channel for the go block
2.        captured-bindings (getThreadBindingFrame) ;; ❷ saving the thread bindings
3.        f (fn state-machine ;; creating a state machine
4.          ([] (aset-all! (AtomicReferenceArray. (int 6)) 0 state-machine 1 1)) ;; ❸ state machine initialization code
5.          ([state]
6.            (let [old-frame (getThreadBindingFrame) ;; ❹ grabbing thread bindings before state machine runs
7.                  ret (try
8.                      (resetThreadBindingFrame (aget-object state 3)) ;; setting thread bindings that were stored in ❷
9.                      (loop [] ;; ❺ core of the state machine
10.                        (let [result
11.                            (case (int (aget-object state 1)) ;; ❻ obtaining the current state and jumping
12.                              1 (take! state 2 ch1)
13.                              2 (let [data (foo (aget-object state 2))]
14.                                  (put! state 3 ch2 data))
15.                                  3 (return-chan state (aget-object state 2))))
16.                            (if (identical? result :recur) (recur) result))) ;; ❼ check if we need to park or not
17.                          (catch Throwable ex ;; ❽ some error handling
18.                            (aset-all! state 2 ex)
19.                            (if (seq (aget-object state 4))
20.                                (aset-all! state 1 (first (aget-object state 4)))
21.                                (throw ex)))
22.                            :recur)
23.                            (finally ;; ❾ restoring thread-locals grabbed in ❹
24.                              (aset-object state 3 (getThreadBindingFrame)) (resetThreadBindingFrame old-frame)))
25.                            (if (identical? ret :recur)
26.                                (recur state)
27.                                ret)))))
28.          state (aset-all! (f) USER-START-IDX c BINDINGS-IDX captured-bindings)] ;; ❿ initialization of the state machine
29.          (run-state-machine-wrapped state)
30.          c)
```

```
(def ch1 (chan))
(def ch2 (chan))
(go
  (let [data (<! ch1])
    (>! ch2 (foo data))))
```

go macro

- Compiles the code into a state machine
 - Using the Inversion Of Control (IOC) strategy
 - Is reminiscent of a classic **generator** pattern with **yield**
- Channel operations are transformed into **callbacks**
- Run by a custom scheduler

```
1. (defn put! [state blk c val]
2.   (if-let [cb (impl/put! c val
3.             (fn-handler
4.               ;; anonymous function, called when someone does take! on the channel c
5.               (fn [ret-val]
6.                 ;; stores the result and the next state in the state machine itself
7.                 (aset-all! state 2 ret-val 1 blk)
8.                 ;; runs the state machine from the callback
9.                 (run-state-machine-wrapped state))))]
10.    (do (aset-all! state 2 @cb 1 blk)
11.        :recur)
12.    nil))
```


Behind the `go` macro: `tools.analyzer`

- A library for:
 - Parsing Clojure code in a host-agnostic way
 - Analysing the resulting AST
- Written in Clojure


Channels (again)

- Actually, not one, but three queues
 - Queue for **put** tasks
 - Queue for **take** tasks
 - Buffer for **data** (optional)

Channels visualized



put into a channel



take from a channel

Timers

- Timeout
 - A special type of channel
 - Automatically closed after set period
- Soft
 - timer resolution is 10ms
- shared
 - Timers created within 10ms from each other use the same object in memory
 - Best not to close them manually, since it could be someone else's timer too

1. (def timer
2. (timeout 1000)) ;; closes after 1 second

Usage example: waiting with timeout


```
1. (def ch (chan))
2.
3. (go
4.   (let [t (timeout 1000)
5.         [data result-chan] (alts! [ch t])] ;; Using "select" on both channels, getting
6.                           ;; the result and the channel it was returned from
7.
8.   (cond (= result-chan t)    ;; if the result channel is the timeout channel,
9.         (println "timeout")  ;; the timeout happened before we could receive data
10.
11.        (= data nil) ;; if we got nil, then someone closed ch (since first check failed)
12.        (println ch "was closed")
13.
14.        :else ;; successfully received data under one second
15.        (println data))))
```

Scheduler

- Platform dependent
 - ClojureScript: goog.async.nextTick
 - Clojure: thread pool

CLJS Scheduler

- Simple stop the world loop
- Only processes so many tasks (1024)
- Re-equeues itself if necessary
- Doesn't run until requested
- Frequency determined by runtime



Do we need a scheduler?

- What scheduler does:
 - Watches over timeout channels
 - Pushes go-threads without external consumers or producers
- If thread pools are available:
 - Decrease reaction time in IO-bound tasks
 - Allows mixing parking and blocking code
- However, scheduler is optional if:
 - Fully reactive approach
 - No threads without external dependencies are allowed
 - (only in single-threaded systems)

Example: required vs optional scheduler

```
1. (go
2.   (<! (timeout 1000))
3.   (println "done"))
```

No external consumers
(scheduler needed, otherwise thread will never complete)

```
1. (def ch (chan))
2.
3. (go
4.   (<! (timeout 1000))
5.   (>! ch "done"))
6.
7. (println (<!! ch))
```

External consumer exists
(no need for scheduler, can work in a reactive manner)

Portability

- Basing go-threads on macros and state machines gives us:
 - Portability across wide range of runtimes
 - No dependencies on platform-threads
 - Ability to manually schedule tasks
- Channels give us:
 - No dependencies on platform-specific async primitives (promises, actors, etc)
 - Clear way to introduce discrete points of separation to our code to create state machines

Problems

;; Clojure

```
1. (def ch (chan))
2.
3. (defn my-async-task []
4.   (let [data (<! ch)] ; compile error: can't use <! outside of go-thread
5.     (do-stuff data)))
6.
7. (go (my-async-task))
```

// Go

```
1. var ch = make(chan int)
2.
3. func myAsyncTask() {
4.   data := <-ch // no problems (backed by runtime)
5.   doStuff(data)
6.
7. func main() { go myAsyncTask() }
```

; Fennel (my port of core.async)

```
1. (local ch (chan))
2.
3. (fn my-async-task []
4.   (let [data (<! ch)] ; no problems (coroutines)
5.     (do-stuff data)))
6.
7. (go (my-async-task))
```

Problems

```
; Clojure (state-machines)

1. (def ch1 (chan))
2. (def ch2 (chan))
3. (def ch3 (chan))
4.
5. (go
6.   ;; compile error:
7.   ;; can't use "syntax" <! as a function
8.   (map <! [ch1 ch2 ch3]))
9.
10. (go
11.  (map
12.    ;; compile error: can't transform lazy loop
13.    ;; into state machine
14.    (fn [c] (<! c)))
15.    [ch1 ch2 ch3]))
```


```
; Fennel (coroutines)

1. (local ch1 (chan))
2. (local ch2 (chan))
3. (local ch3 (chan))
4.
5. (go
6.   ;; OK:
7.   ;; <! is a function, we're good
8.   (map <! [ch1 ch2 ch3]))
9.
10. (go
11.  (map
12.    ;; Also OK:
13.    ;; can yield across function calls
14.    (fn [c] (<! c)))
15.    [ch1 ch2 ch3]))
```


Conclusion

- Pros:
 - Using macros as means for thread building allows great portability
 - State-machines allow for platform-independent way to start and resume code
 - Channels allow us to specify where parking happens
- Cons:
 - Nuances and unexpected behavior due to the lack of platform support
 - Wrappers may be required if we want to interface with other asynchronous solutions (virtual threads, callbacks)

Thanks!



slides



team.health-samurai.io